



Final answers on identical questions · ordinary and recurrent



All models use the same 345,600 graph/count queries. Ordinary Qwen uses 3-shot chat generation; recurrent models use raw prompts/restricted readout.

Left: recurrent readout at N. Right: actual returned letter; correct letters at wrong loops count here, so this is NOT answer + exact-stop success.

No trajectory or loop-timing metric is assigned to ordinary generation. CE has no learned stop and is absent from the right panel.