





Ordinary Qwen · current full graph/count benchmark



Same 1,350 graphs, starts, rule order and requests 1-256 as the recurrent suite; 345,600 unconstrained generations. Existing 3-shot chat prompt; greedy, max 8 new tokens. Invalid responses are failures. Prompt/readout differ from recurrent inference. Pointwise intervals use 1,350 graphs; requests share graph clusters. Generation tokens are not recurrent transitions.