


Ordinary base checkpoint · 6% final-answer accuracy



Original seed-17 test: 1,000 questions, 125 per depth, greedy 3-shot generation; 95% Wilson intervals.
Historical context: this is a different dataset/prompt/inference procedure from the recurrent population comparison.